

Automatic Music Accompanist

Anyi Rao

Nanjing University
anyirao@smail.nju.edu.cn

Francis Lau

The University of Hong Kong
fcmlau@cs.hku.hk

ABSTRACT

Automatic musical accompaniment is where a human musician is accompanied by a computer musician. The computer musician is able to produce musical accompaniment that relates musically to the human performance. The accompaniment should follow the performance using observations of the notes they are playing.

This paper describes a complete and detailed construction of a score following and accompanying system using Hidden Markov Models (HMMs). It details how to train a score HMM, how to deal with polyphonic input, how this HMM work when following score, how to build up a musical accompanist. It propose a new parallel hidden markov model for score following and a fast decoding algorithm to deal with performance errors.

1. INTRODUCTION

Accompanists may not always be available when needed, or available accompanists may not have sufficient technical ability to provide adequate accompaniment. A solution for many musicians is to make use of recorded or computer-generated accompaniment where the accompaniment is static, i.e. never changing from one performance to another. This forces the musician to adapt their playing to synchronize with the accompaniment. It is more natural for the musician, though, if the accompaniment adapts to the performer, particularly as a musicians playing tends to be 'free'.

To dynamically synchronize the accompaniment with the performance by the musician, the accompanist should track the performers progress through the score of the piece as they play. Score following is the process whereby a musician follows another musicians playing of a musical piece, by tracking their progress through the score of that piece. The term is most commonly used in the context of computer-generated accompaniment, where one or more of the musicians involved are artificial rather than human. The purpose of the research outlined in this paper is to construct a automatic accompaniment.

In live performance, score following must be on-line real-time, i.e. producing accompaniment in time with the soloists playing. This places extra challenges for the score follower. The system has a more limited amount of information available for analysis: only the notes that have been played so far, as opposed to having the whole performance

to analyse. It requires fast computation speed. The accompanist needs finish one accompaniment before the next note comes.

Our contributions are as follows,

- 1 Our work is the first free open-source Windows based automatic music follower and accompanist to our best knowledge.
- 2 We construct a comprehensive system and show how it works with detailed theoretical induction, including score follower training/decoding and score accompanist.
- 3 We propose a fast decoding algorithm, reduced computational complexity from $O(n^2)$ down to $O(n)$. It is able to work in real time with practical length scores.
- 4 We build up two hands parallel HMM to improve accuracy and computational speed.

Background

There are several reasons why a musician may not perform the piece exactly as written. Changes may be added by mistake: 1. A wrong note is played; 2. Extra notes are added; 3. Scored notes are missed out; 4. The musician loses their place in the music or starts playing from the wrong point in the score. 5. The tempo speeds up or slows down unintentionally. Also changes may be added deliberately, as the musician adds their own interpretations to the music: 1. The musician adds embellishments such as trills, to decorate the notes; 2. The tempo speeds up or slows down deliberately, for musical effect; 3. The piece being played may have rubato or free/improvised sections, where the musician is free to vary the tempo and notes played according to their own choice.

Terms

Performance: In the context of this project, a performance is defined specifically as the situation where a solo musician (soloist), such as a flute player or singer, performs a piece of music. The solo musician would be accompanied by another musician (accompanist) on an instrument such as piano. This may be in a concert or similar scenario, performing to an audience, but this condition is not mandatory. What is important is that the soloist is making an attempt to play through the piece in a linear fashion, from start to finish.

Performer/Soloist: The solo musician who is performing the piece; what they play is the most important part of the performance for any audience that may be listening.

Accompanist: The musician who is playing the accompaniment; supporting the soloists performance.

Melody/Solo melody: The music that is being played by the soloist.

Accompaniment: The music which is played by an accompanist, during the performance of the soloist. Accompaniment can be thought of background music which is designed to enhance what the soloist is playing and support the soloists performance.

Score follower: A computer accompanist that follows the solo melody through the score as it is being played, to produce accompaniment relative to where the soloist is in the score.

2. HIDDEN MARKOV MODELS

A musical score is divided up into a sequence of musical events. (for example one note or one beat can be considered as one modellable musical event)

The score follower uses a Hidden Markov Model to represent these musical events, and uses a decoding algorithm to estimate what state the performer is most likely to be in at that time, i.e. which musical event in the score the performer is currently playing.

2.1 HMM Structure

We define the observation states as 12 notes in the western musical chromatic scale. We ignore octave differences between notes and merely consider 12 possible observations: $\{ C, C\sharp, D, E\flat, E, F, F\sharp, G, G\sharp, A, B\flat, B \}$, as shown in Figure 1. The hidden states base on beat and encode the information relates to the beat as detailed in the below section 2.1.1. The paraments λ of the HMM contains three parts $\{ \pi, A, B \}$ denoted below.

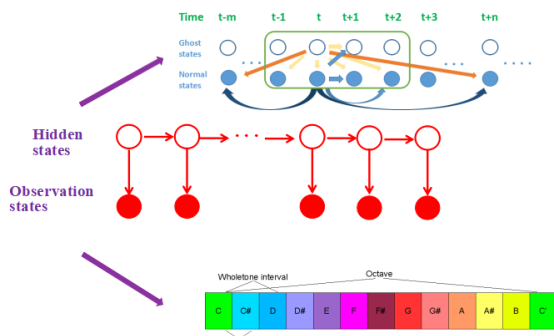


Figure 1. Hidden Markov Model structure.

- N : the number of possible states. We use N symbols S_1, S_2, \dots, S_N to denote them.
- M : the number of possible observations.
- π : the prior (initial) state distribution.
 $\pi = (\pi_1, \pi_2, \dots, \pi_N)$ and $\pi_i = \Pr(Q_1 = S_i)$.

- A : the state transition matrix. $A_{ij} = \Pr(Q_t = S_j | Q_{t-1} = S_i)$ is the probability of the next state being S_j if the current one is S_i , $1 \leq i, j \leq N$. Note that A does not change when t changes.
- B : the observation probability matrix. Instead of denoting one probability as B_{jk} , we use $b_j(k) = \Pr(O_t = V_k | Q_t = S_j)$ to denote the probability of the observation being V_k when the state is S_j , $1 \leq j \leq N$, $1 \leq k \leq M$. And, B does not change when t changes.

2.1.1 beat-based representation

If there is a simple tune for which each note is of the same length, the naive choice is to model each note as an individual HMM state.

But when music pieces become more complex, it is no longer realistic to model each note as a new state, and instead the more pertinent aspect to model as a state is each beat, or a fraction of each beat. For such cases, it was necessary to consider how the timing information within the score should be modelled (in addition to how the notes should be modelled).

The two obvious ways to model a note that is held for longer than one state (i.e. notes that extend over a beat or more) are:

1. Allow states with self-transitions, so the HMM stays in a given state while a note is being held and only moves out of that state when the note is released.
2. Have a finite number of states representing each note that is longer than one state, proportional to the length of the note (for example if each state represents one beat and a note is three beats long, represent it as three sequential states).

The more successful option here is the second [1] with more flexibility to vary the accompaniment and it also able to encode notes of different lengths into the HMM.

2.1.2 Errors representation

There are three classes of probable errors [2]:

- **WRONG:** An incorrect note is played in place of the correct note.
- **SKIP:** A note in the score is missed out altogether.
- **EXTRA:** An extra, unscored note is added in the performance.

The Hidden Markov Model processes such errors by the soloist, as they happen, by taking a specific path through the normal and ghost states. The paths for each class of error are shown in Figure 2.

2.2 Training

We train the HMM involving getting the maximum probability of being in the correct normal state or ghost state, given a sequence of observations. We define four variables.

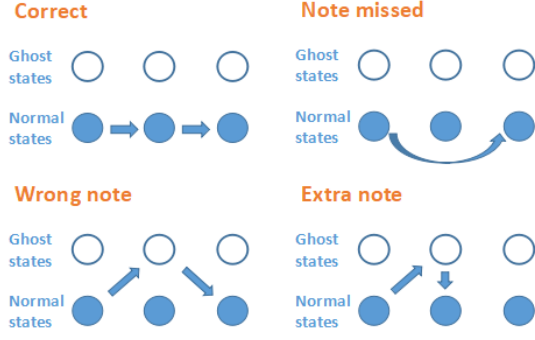


Figure 2. Typical deviations from a score and the HMM hidden state transitions associated with these deviations.




Figure 3. All allowed transitions from the first normal/ghost state pair.

- $\alpha_t(i)$: $\alpha_t(i) = Pr(o_{1:t}, Q_t = S_i | \lambda)$ with recursion: $\alpha_{t+1}(i) = (\sum_{j=1}^N \alpha_t(j) A_{ji}) b_j(o_{t+1})$
- $\beta_t(i)$: $\beta_t(i) = Pr(o_{t+1:T} | Q_t = S_i, \lambda)$ with calculation: $\beta_t(i) = \sum_{j=1}^N A_{ij} b_j(o_{t+1}) \beta_{t+1}(j)$
- $\gamma_t(i)$: $\gamma_t(i) = Pr(Q_t = S_i | o_{1:T}, \lambda)$
- $\xi_t(i, j) = Pr(Q_t = S_i, Q_{t+1} = S_j | o_{1:T}, \lambda)$.

The ξ variable involves three other values: t (the time) and (i, j) which are state indexes. Comparing the definition of γ and ξ , we immediately get (by the law of total probability):

$$\gamma_t(i) = \sum_{j=1}^N \xi_t(i, j) \quad (1)$$

The parameters $\lambda = (\pi, A, B)$ can be updated using γ and ξ . Using the definition of conditional probabilities, we have

$$\xi_t(i, j) Pr(o_{1:T} | \lambda) = Pr(Q_t = S_i, Q_{t+1} = S_j, o_{1:T} | \lambda). \quad (2)$$

we can find the probability $Pr(Q_t = S_i, Q_{t+1} = S_j, o_{1:T} | \lambda)$ and use it to compute $\xi_t(i, j)$. This probability can be factored into the product of four probabilities: $\alpha_t(i)$, A_{ij} , $b_j(o_{t+1})$ and $\beta_{t+1}(j)$

$$\xi_t(i, j) = \frac{\alpha_t(i) A_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{Pr(o_{1:T} | \lambda)} \quad (3)$$

The entire training algorithm are shown in Algorithm 1.

Algorithm 1 Training Algorithm

- 1: Initialize the parameters $\lambda^{(1)}$ (e.g., randomly)
- 2: $\tau \leftarrow 1$
- 3: **while** the likelihood has not converged **do**
- 4: Use the forward procedure to compute $\alpha_t(i)$ for all t ($1 \leq t \leq T$) and all i ($1 \leq i \leq N$) based on $\lambda^{(\tau)}$
- 5: Use the backward procedure to compute $\beta_t(i)$ for all t ($1 \leq t \leq T$) and all i ($1 \leq i \leq N$) based on $\lambda^{(\tau)}$
- 6: Compute $\gamma_t(i)$ for all t ($1 \leq t \leq T$) and all i ($1 \leq i \leq N$) according to the equation in Table 1
- 7: Compute $\xi_t(i, j)$ for all t ($1 \leq t \leq T-1$) and all i, j ($1 \leq i, j \leq N$) according to the equation in Table 1
- 8: Update the parameters to $\lambda^{(\tau+1)}$

$$\pi_i^{(\tau+1)} = \gamma_1(i) \quad (4)$$

$$A_{ij}^{(\tau+1)} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)} \quad (5)$$

$$b_j^{(\tau+1)}(k) = \frac{\sum_{t=1}^T \mathbb{1}_{o_t = k} \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)} \quad (6)$$

- 9: $\tau \leftarrow \tau + 1$

10: **end while**

2.3 Real time decoding

The aim is to find the most probable hidden state sequence that could generate the observations sequence produced by hearing the soloists playing. In the score followers developed during this project, a revised Viterbi algorithm is used to find out which state the soloist is most likely to be in (given the sequence of observations of what notes the soloist has most recently played).

Implemented in the traditional fashion [3], this algorithm finds the globally optimum path through the Hidden Markov Model states to the most probable current state, using the history of observations seen. But this causes huge computational complexity and the system cannot be used in practice¹.

Although one might consider some pruning techniques to reduce computational complexity, pruning is not valid within the context of handling arbitrary skips since skips rarely occur compared to other state transitions. Therefore, it seems necessary to introduce some constraints to the performance HMM.

The problem with large computational complexity arises from the non-zero values of the transition probability a_{ij} for large $|i - j|$. Here, we assume the transition probability can be summarised as

$$a_{ij} = \tilde{A}_{ij} + \mu. \quad (7)$$

where \tilde{A}_{ij} is a band matrix satisfying $\tilde{A}_{ij} = A_{ij}$ when $i - W_1 \leq j \leq i + W_2$, otherwise $\tilde{A}_{ij} = 0$, which describes transitions within neighbouring states. μ is a prior

¹ Most classical musical pieces have $O(100 - 10000)$ chords. For example, the solo piano part of Rachmaninoffs piano concerto No. 3 d-moll has $N \simeq 5000$ chords only in the first movement.

distribution got from training part [4] depicting arbitrary repeats/skips².

$$a_{ij} = \mu, \quad \text{for } j < i - W_1 \quad \text{or} \quad j > i + W_2 \quad (8)$$

where W_1 and W_2 are small positive integers which define a neighbourhood of states.

Given an observation sequence $o_{1:T}$, we use a new variable δ , defined by Equation (9), to find the best path. W is the sliding window width and $W = W_1 + W_2 + 1$. δ has recursive relationship, as shown in equation (10)

$$\delta_t(i) = \max_{Q_{1:t-1}} Pr(Q_{1:t-1}, o_{1:t}, Q_t = S_i | \lambda) \quad (9)$$

$$\delta_{t+1}(i) = \max_{1 \leq j \leq N} (\delta_t(j) A_{ji} b_i(o_{t+1})) \quad (10)$$

Algorithm 2 Decoding Algorithm

- 1: **Initialization:** $\delta_1(i) = \pi_i b_i(o_1), \psi_1(i) = 0$ for all $1 \leq i \leq N$
- 2: **Slide algorithm window:**
- 3: **Forward recursion:** For $t = 1, 2, \dots, T - 2, T - 1$ and all $1 \leq i \leq N$

$$\delta_{t+1}(i) = \max(\delta_t(j) a_{ji} b_i(o_{t+1})) \quad (11)$$

$$\psi_{t+1}(i) = \operatorname{argmax}(\delta_t(j) a_{ji}) \quad (12)$$

- 4: **Output:** The optimal state q_T is determined by

$$q_T = \operatorname{arg} \max_{1 \leq i \leq N} \delta_{t+1}(i) \quad (13)$$

Theoretical evaluation

We can rewrite equation (11) as below:

$$\delta_{t+1}(i) = b_i(o_{t+1}) \max \left\{ \max_{j \in \text{nbh}(i)} [\delta_t(j) A_{ji}], \max_j [\delta_t(j) \mu] \right\} \quad (14)$$

where $\text{nbh}(i) = \{j | j - W_1 \leq i \leq j + W_2\}$ denotes the set of neighbouring states of i . Since the factor $\max_j [\delta_t(j) \mu]$ in the last equation is independent of i and can be calculated with $O(N)$ complexity, the decoding algorithm expression has $O(WN)$ computation complexity compared with previous $O(N^2)$ complexity. Therefore, a fast Viterbi algorithm can be used efficiently for the HMM if $W \ll N$.

2.4 Two hands parallel HMM

We construct a two hands parallel HMM, with each hand as part HMMs corresponding to the HMM described above. The two then merged their outputs, assuming there is no hand crossing in performance.

² performers are likely to resume their performance from the beginning of a sentence/phrase when they make mistakes [5]

The two part HMMs transits and outputs an observed symbol at each time. The state space of the parallel HMM is given as a triplet $k = (\eta, f_L, f_R)$ of the hand information, where η indicate which of the HMMs works, and f_L and f_R indicate the current states of the part HMMs. [6]

3. ACCOMPANIST

A soloist will naturally incorporate expressive features in the playing, involving shaping of the tempo and intensity of the playing in ways not explicitly represented in the score.

A human accompanist would not wait for every note to be played by the soloist before playing accompaniment. Instead they anticipate that the soloist will move onto the next note in the score and play the appropriate accompaniment, then use the incoming information from the soloist to update their belief of where the soloist is in the score and adjust their accompaniment if necessary.

In a similar fashion, this system uses the Hidden Markov Model representation to work out what the next sequential state is, playing the accompaniment for that state at the time it expects the next state to occur. As it receives and processes the soloists actual input and locates the HMM state that the soloist has actually reached, it adjusts the accompaniment if necessary.

3.1 Beat tracking

We implement beat tracking to monitor the performance's tempo and provide a reference to the accompaniment speed.

The system incorporates a simple version of beat tracking. This allows small tempo fluctuations to be tracked, and the soloists output to be anticipated in a timely fashion. Modelling the score by temporal units assisted us greatly with including beat tracking in the accompanist. Our implementation was simpler than [7] but was effective.

The accompanist used an internal tempo measure that was continually adjusted to match the soloists estimated current tempo, using a local window of notes recently played by the soloist and measuring the time in between those notes (relative to the notes expected durations). If the soloist is currently judged to be in a ghost state (i.e. they have deviated from the score), then the last input is not considered as valid for use in updating the tempo. If, though, the soloist is currently judged to be in a normal state (i.e. they can be found on the score), then the score follower works out how long the previous note should have been and compares this with the actual length of the last note. The current tempo is based on an average of the recent (valid) tempo observations. The largest and smallest tempo observations are ignored and a mean is taken of the remaining tempo observations, to generate an estimate of the current tempo.

3.2 Controlling dynamics of the performance

The system can track the volume of the soloists playing using MIDI information and replicate that volume in the dynamic level of the accompaniment output, playing the accompaniment at a very slightly lower volume than the

soloist. In this way the system allows the soloist line to be prominent but also matches the dynamic markings of their playing. We felt it was more important to be responsive to the soloists dynamic interpretations than to allow the accompanist to play at a dynamic marking independent of the soloists dynamics.

3.3 Rule-based Reactive accompanist

If a human accompanist hears their soloist deviate slightly from the score, it takes time for the accompanist to relocate the soloist and adjust their playing from the expected accompaniment to the accompaniment matching the soloist.

It would be reasonable to have the computer accompanist only respond to a deviation on the next state after a deviation from the score was identified: replicating the slight delay that a human accompanist would also have. This is on the assumption that the states are modelled such that they are close enough together in timing for the delay not to be too noticeable.

We use a musical accompanying rule to generate our accompaniment [8], i.e. a chord match with some certain chords.

Acknowledgments

Thanks to the support of the University of Hong Kong Computer Science Department summer research internship programme.

4. REFERENCES

- [1] A. Jordanous and A. Smaill, "Investigating the role of score following in automatic musical accompaniment," in *J. New Music Research*, 2009, pp. 197–209.
- [2] N. Orio and F. Déchelle, "Score following using spectral analysis and hidden markov models," in *ICMC: International Computer Music Conference*, 2001, pp. 1–1.
- [3] L. R. Rabiner, "A tutorial on hidden markov models and selected applications in speech recognition," *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257–286, 1989.
- [4] B. Couasnon and J. Camillerapp, "Using grammars to segment and recognize music scores," in *International Association for Pattern Recognition Workshop on Document Analysis Systems*, 1994, pp. 15–27.
- [5] T. Nakamura, E. Nakamura, and S. Sagayama, "Real-time audio-to-score alignment of music performances containing errors and arbitrary repeats and skips," *IEEE/ACM Transactions on Audio, Speech and Language Processing (TASLP)*, vol. 24, no. 2, pp. 329–339, 2016.
- [6] E. Nakamura, N. Ono, Y. Saito, and S. Sagayama, "Merged-output hidden markov model for score following of midi performance with ornaments, desynchronized voices, repeats and skips," *algorithms*, vol. 21, p. 8, 2014.
- [7] S. Dixon, "Automatic extraction of tempo and beat from expressive performances," *Journal of New Music Research*, vol. 30, no. 1, pp. 39–58, 2001.
- [8] A. Friberg, "Generative rules for music performance: A formal description of a rule system," *Computer Music Journal*, vol. 15, no. 2, pp. 56–71, 1991.